

Diffusion models

Цель

- Познакомится с концептом диффузионных моделей, рассмотреть различные практические применения и расширения диффузионных моделей

План


- Прямой и обратный диффузионные процессы (DDPM)
- Условная генерация с помощью диффузионных моделей
- Latent Diffusion Models и Stable Diffusion
- Применения диффузионных моделей
- Методы дообучения диффузионных моделей
- Методы добавления новых условий в обученные модели
- Инструменты для работы с диффузионными моделями

Прямой и обратный диффузионные процессы (DDPM)


Denoising Diffusion Probabilistic Models

[Denoising Diffusion Probabilistic Models, Arxiv](#)


Прямой и обратный диффузионные процессы



Прямой и обратный диффузионные процессы



Прямой диффузионный процесс




- В теории:

$$x_0 \Rightarrow x_1 \Rightarrow \dots \Rightarrow x_T$$
$$q(x_t|x_{t-1}) = \mathbb{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t \mathbb{I}); q(x_T|x_0) = \mathbb{N}(x_T; 0, 1)$$
$$\beta_0 = 0, \beta_T = 1, \beta_t \geq \beta_{t-1}, T = 1000$$

$$q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1})$$


Прямой диффузионный процесс



- На практике:

$$\epsilon \sim \mathcal{N}(0,1)$$
$$x_t = \sqrt{1 - \beta_t} * x_{t-1} + \sqrt{\beta_t} * \epsilon$$

Прямой диффузионный процесс



- На практике:


$$\epsilon \sim N(0,1)$$

$$\text{вводя } \alpha_t = 1 - \beta_t, \bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$
$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} * \epsilon$$

β_t выбираются так, чтобы $\bar{\alpha}_t \rightarrow 0$ и $q(x_T|x_0) \approx N(x_T; 0,1)$

Обратный диффузионный процесс

Diffused Data Distributions




Можем ли мы
аппроксимировать $q(x_{t-1}|x_t)$?

Да, мы можем использовать
нормальное распределение,
при условии, что изменение β_t
малое на каждом прямом
диффузионном шаге.

На самом деле описать
 $q(x_{t-1}|x_t, x_0)$ достаточно
легко, об этом далее.

Обратный диффузионный процесс



Обратный диффузионный процесс

Вариационная нижняя оценка

$$\mathbb{E}_{q(\mathbf{x}_0)} [-\log p_\theta(\mathbf{x}_0)] \leq \mathbb{E}_{q(\mathbf{x}_0)q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \left[-\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] =: L$$

$$L = \mathbb{E}_q \left[\underbrace{D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0)||p(\mathbf{x}_T))}_{L_T} + \sum_{t>1} \underbrace{D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)||p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t))}_{L_{t-1}} \underbrace{- \log p_\theta(\mathbf{x}_0|\mathbf{x}_1)}_{L_0} \right]$$

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}),$$

$$\text{where } \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{1 - \beta_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}\mathbf{x}_t \text{ and } \tilde{\beta}_t := \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t}\beta_t$$

Обратный диффузионный процесс

Вариационная нижняя оценка

- Поскольку мы предполагаем, что $q(x_{t-1}|x_t, x_0)$ и $p_\theta(x_{t-1}|x_t)$ - нормальные распределения:

$$L_{t-1} = D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)||p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)) = \mathbb{E}_q \left[\frac{1}{2\sigma_t^2} ||\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) - \mu_\theta(\mathbf{x}_t, t)||^2 \right] + C$$

$$\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) = \frac{1}{\sqrt{1 - \beta_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon \right) \quad \mu_\theta(\mathbf{x}_t, t) = \frac{1}{\sqrt{1 - \beta_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right)$$

$$L_{t-1} = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}_0), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[\frac{\beta_t^2}{2\sigma_t^2(1 - \beta_t)(1 - \bar{\alpha}_t)} ||\epsilon - \underbrace{\epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)}_{\mathbf{x}_t} ||^2 \right] + C$$

Обратный диффузионный процесс

Вариационная нижняя оценка

$$L_{t-1} = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}_0), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})} \left[\underbrace{\frac{\beta_t^2}{2\sigma_t^2(1 - \beta_t)(1 - \bar{\alpha}_t)}}_{\lambda_t} \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2 \right]$$

$$L_{\text{simple}} = \mathbb{E}_{\mathbf{x}_0 \sim q(\mathbf{x}_0), \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), t \sim \mathcal{U}(1, T)} \left[\|\epsilon - \epsilon_\theta(\underbrace{\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon}_{\mathbf{x}_t}, t)\|^2 \right]$$

Расписание шума

Косинусное и линейное расписания

$$\bar{\alpha}_t = \frac{f(t)}{f(0)}, f(t) = \cos\left(\left(\frac{t}{T} + s\right) * \frac{\pi}{2}\right)^2$$

$$\beta_t = 1 - \frac{\bar{\alpha}_t}{\bar{\alpha}_{t-1}}$$




Figure 5. $\bar{\alpha}_t$ throughout diffusion in the linear schedule and our proposed cosine schedule.

Расписание шума

Косинусное и линейное расписания



Figure 3. Latent samples from linear (top) and cosine (bottom) schedules respectively at linearly spaced values of t from 0 to T . The latents in the last quarter of the linear schedule are almost purely noise, whereas the cosine schedule adds noise more slowly

Процесс обучения и применения диффузионной модели

Algorithm 1 Training


```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$   
6: until converged
```

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

Условная генерация с
помощью диффузионных
моделей

Генерация без условия



Условная генерация

- Очевидно, что мы можем переучить модель с нуля, добавив в архитектуру механизмы, которые смогут получать и обрабатывать дополнительный сигнал, например эмбединг текста или картинки. Удобнее всего использовать cross-attention и блоки модели-трансформера
- Обучение новой модели – долгий и дорогостоящий процесс, можно ли обойтись без переобучения?

Classifier-guidance

- Пусть $p_\phi(y|x_t)$ - классификатор, обученный на зашумленных объектах, можем использовать его градиенты в качестве поправки к оценке предсказанного нашей моделью шума

Algorithm 2 Classifier guided DDIM sampling, given a diffusion model $\epsilon_\theta(x_t)$, classifier $p_\phi(y|x_t)$, and gradient scale s .

Input: class label y , gradient scale s
 $x_T \leftarrow$ sample from $\mathcal{N}(0, \mathbf{I})$
for all t from T to 1 **do**
 $\hat{\epsilon} \leftarrow \epsilon_\theta(x_t) - \sqrt{1 - \bar{\alpha}_t} \nabla_{x_t} \log p_\phi(y|x_t)$
 $x_{t-1} \leftarrow \sqrt{\bar{\alpha}_{t-1}} \left(\frac{x_t - \sqrt{1 - \bar{\alpha}_t} \hat{\epsilon}}{\sqrt{\bar{\alpha}_t}} \right) + \sqrt{1 - \bar{\alpha}_{t-1}} \hat{\epsilon}$
end for
return x_0

Classifier-guidance



Figure 3: Samples from an unconditional diffusion model with classifier guidance to condition on the class "Pembroke Welsh corgi". Using classifier scale 1.0 (left; FID: 33.0) does not produce convincing samples in this class, whereas classifier scale 10.0 (right; FID: 12.0) produces much more class-consistent images.

Classifier-free guidance

- Переобучим модель и добавим в качестве дополнительного входа в сеть, предсказывающую шум, метку класса / эмбединг текста или картинки $\epsilon = \epsilon(x_t, t, y)$. Также будем с некоторой вероятностью подавать на вход пустое условие, чтобы модель могла генерировать без условия $\epsilon = \epsilon(x_t, t)$

$$\begin{aligned}\nabla_{\mathbf{x}_t} \log p(y|\mathbf{x}_t) &= \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t|y) - \nabla_{\mathbf{x}_t} \log p(\mathbf{x}_t) \\ &= -\frac{1}{\sqrt{1 - \bar{\alpha}_t}} \left(\epsilon_{\theta}(\mathbf{x}_t, t, y) - \epsilon_{\theta}(\mathbf{x}_t, t) \right)\end{aligned}$$

$$\begin{aligned}\bar{\epsilon}_{\theta}(\mathbf{x}_t, t, y) &= \epsilon_{\theta}(\mathbf{x}_t, t, y) - \sqrt{1 - \bar{\alpha}_t} w \nabla_{\mathbf{x}_t} \log p(y|\mathbf{x}_t) \\ &= \epsilon_{\theta}(\mathbf{x}_t, t, y) + w(\epsilon_{\theta}(\mathbf{x}_t, t, y) - \epsilon_{\theta}(\mathbf{x}_t, t)) \\ &= (w + 1)\epsilon_{\theta}(\mathbf{x}_t, t, y) - w\epsilon_{\theta}(\mathbf{x}_t, t)\end{aligned}$$

Classifier-free guidance



Non-guidance




$\omega = 1$




$\omega = 3$

Latent Diffusion Models и Stable Diffusion


Latent Diffusion Models (LDM)



Stable Diffusion (1.5)




Stable Diffusion (1.5)




Stable Diffusion (1.5)

CrossAttnDownBlock2D


CrossAttnDownBlock2D



ResnetBlock2D




Transformer2DModel




Stable Diffusion (1.5)

UNetMidBlock2DCrossAttn


UNetMidBlock2DCrossAttn



ResnetBlock2D




Transformer2DModel




Stable Diffusion (1.5)

CrossAttnUpBlock2D



Применения диффузионных моделей

Text-to-image



vibrant portrait painting of Salvador Dalí with a robotic half face



a shiba inu wearing a beret and black turtleneck



a close up of a handpalm with leaves growing from it



an espresso machine that makes coffee from human souls, artstation



panda mad scientist mixing sparkling chemicals, artstation



a corgi's head depicted as an explosion of a nebula

[Hierarchical Text-Conditional Image Generation with CLIP Latents, Arxiv](#)

Image Inpainting



Image Inpainting






Figure 2. **Overview of our approach.** RePaint modifies the standard denoising process in order to condition on the given image content. In each step, we sample the known region (*top*) from the input and the inpainted part from the DDPM output (*bottom*).

Super-resolution & Colorization




Super-resolution



Colorization

Music Generation (Riffusion)



Методы дообучения диффузионных моделей

Дообучение диффузионных моделей

Стиль




Textual Inversion
Dreambooth

Объект




Textual Inversion
Dreambooth

Условие



Retraining
ControlNet
IP-Adapter

Textual Inversion




[An Image is Worth One Word: Personalizing Text-to-Image Generation using Textual Inversion, Arxiv](#)

Textual Inversion



[An Image is Worth One Word: Personalizing Text-to-Image Generation using Textual Inversion, Arxiv](#)

Dreambooth



Dreambooth



[DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation, Arxiv](#)


Dreambooth with LoRA

Снижаем требования к VRAM за счет обучения добавок $\Delta W = BA$ к слоям модели


$$h = W_0 x + \Delta W x = W_0 x + BAx$$

$$W \in R^{N \times M}, B \in R^{N \times r}, A \in R^{r \times M}$$

Получаем адаптер к весам базовой модели.



civitai.com



civitai.com

Смешивание LoRA



No LoRA



Add detail LoRA



Van Gogh LoRA



Add detail + Van Gogh LoRA

Методы добавления новых условий в обученные модели

ControlNet




Figure 7: Controlling Stable Diffusion with various conditions **without prompts**. The top row is input conditions, while all other rows are outputs. We use the empty string as input prompts. All models are trained with general-domain data. The model has to recognize semantic contents in the input condition images to generate images.

ControlNet




Figure 2: A neural block takes a feature map x as input and outputs another feature map y , as shown in (a). To add a ControlNet to such a block we lock the original block and create a trainable copy and connect them together using zero convolution layers, *i.e.*, 1×1 convolution with both weight and bias initialized to zero. Here c is a conditioning vector that we wish to add to the network, as shown in (b).




Figure 3: Stable Diffusion's U-net architecture connected with a ControlNet on the encoder blocks and middle block. The locked, gray blocks show the structure of Stable Diffusion V1.5 (or V2.1, as they use the same U-net architecture). The trainable blue blocks and the white zero convolution layers are added to build a ControlNet.

ControlNet




Figure 4: The sudden convergence phenomenon. Due to the zero convolutions, ControlNet always predicts high-quality images during the entire training. At a certain step in the training process (*e.g.*, the 6133 steps marked in bold), the model suddenly learns to follow the input condition.




Figure 6: Composition of multiple conditions. We present the application to use depth and pose simultaneously.

T2I-Adapter




Figure 3. The overall architecture is composed of two parts: 1) a pre-trained stable diffusion model with fixed parameters; 2) several T2I-Adapters trained to align internal knowledge in T2I models and external control signals. Different adapters can be composed by directly adding with adjustable weight ω . The detailed architecture of T2I-Adapter is shown in the lower right corner.

IP-Adapter

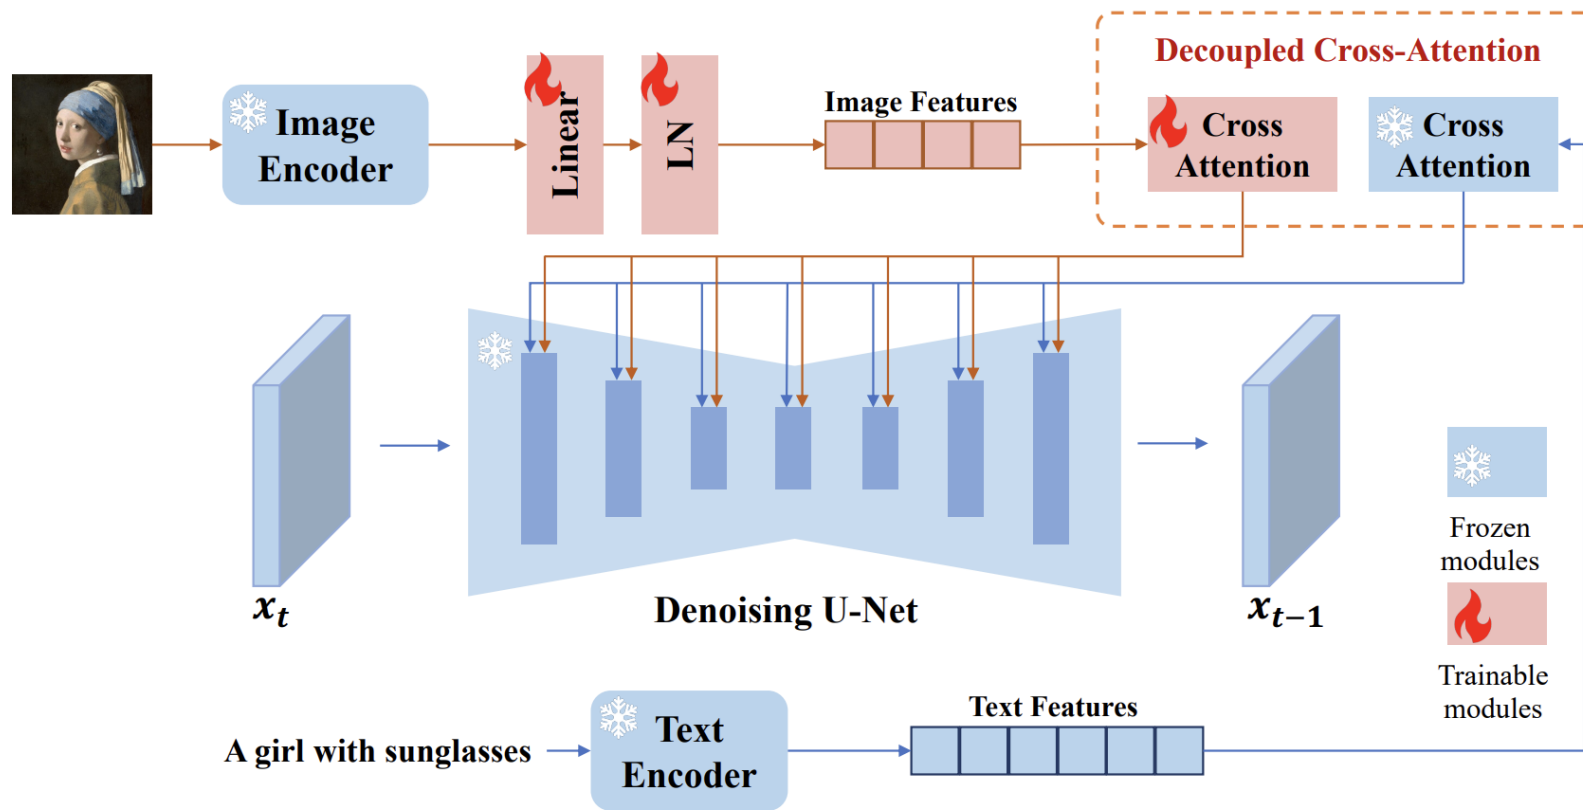


Figure 2: The overall architecture of our proposed IP-Adapter with decoupled cross-attention strategy. Only the newly added modules (in red color) are trained while the pretrained text-to-image model is frozen.

IP-Adapter





Figure 4: The generated images of different diffusion models with our proposed IP-Adapter. The IP-Adapter is only trained once.

IP-Adapter-FaceID-Portrait



<https://huggingface.co/h94/IP-Adapter-FaceID>

Собираем все вместе




Figure 5: Visualization of generated samples with image prompt and additional structural conditions. Note that we don't need fine-tune the IP-Adapter.

Инструменты для работы с диффузионными моделями

AUTOMATIC1111/stable-diffusion-webui

Stable Diffusion checkpoint
protogenX34OfficialR_1.ckpt [60fe2f34]

txt2img img2img Extras PNG Info Checkpoint Merger Train Tokenizer Settings Extensions

green sapling rowing out of ground, mud, dirt, grass, high quality, photorealistic, sharp focus, depth of field

Negative prompt (press Ctrl+Enter or Alt+Enter to generate)

26/75

Generate

Style 1: None Style 2: None

Sampling method: Euler a Sampling steps: 20

☐ Restore faces ☐ Tiling ☐ Hires. fix

Width: 512 Height: 512 Batch count: 4 Batch size: 1

CFG Scale: 12

Seed: 1441787169

Script: None

26/75

green sapling rowing out of ground, mud, dirt, grass, high quality, photorealistic, sharp focus, depth of field

Steps: 20, Sampler: Euler a, CFG scale: 12, Seed: 1441787169, Size: 512x512, Model hash: 60fe2f34, Model: protogenX34OfficialR_1


Time taken: 8.62s Torch active/reserved: 3699/4702 MiB, Sys VRAM: 7020/24576 MiB (28.56%)

API • Github • Gradio • Reload UI

python: 3.10.6 • torch: 1.13.1+cu117 • xformers: N/A • gradio: 3.15.0 • commit: 983167e6

<https://github.com/AUTOMATIC1111/stable-diffusion-webui>

ComfyUI



<https://github.com/comfyanonymous/ComfyUI>

Заключение

- Рассмотрели основной концепт диффузионных моделей
- Поговорили об условной генерации и разных применениях диффузионных моделей
- Затронули тему дообучения моделей и добавления новых условий в уже существующие модели

Полезные ссылки

- [Denoising Diffusion Implicit Models для ускорения сэмплирования](#)
- [Tutorial по диффузионным моделям на CVPR 2022](#) (взгляд на диффузионные модели с точки зрения дифференциальных уравнений)
- [Блог Lillian Weng о диффузионных моделях](#)
- [Статья о Stable Diffusion 3 \(DiT архитектура\)](#)